

Lecture #3: Online Learning*

Janardhan Rao (Jana) Doppa

School of EECS, Washington State University

* Slides partly based on Koby Crammer

Formal setting – Classification

- **Instances**

$$\mathbf{x} \in \mathcal{X}$$

- emails

- **Labels**

$$y \in \mathcal{Y} = \{-1 ; 1\}$$

- Spam vs. non-spam

- **Prediction rule**

$$f(\mathbf{x}) = \hat{y}$$

- Linear prediction rule

- **Loss**

$$\ell(\hat{y}, y) \in \mathbb{R}_+$$

- No. of mistakes

Predictions

- Continuous predictions : $f : \mathcal{X} \rightarrow \mathbb{R}$
 - Label $\text{sign}(f(\mathbf{x}))$
 - Confidence $|f(\mathbf{x})|$
- Linear Classifiers
 - Prediction :
$$\begin{aligned}\hat{y} &= \text{sign}(f(\mathbf{x})) \\ &= \arg \max_{y \in \mathcal{Y}} \mathbf{w} \cdot \Phi(\mathbf{x}, y) \\ &= \text{sign}(\mathbf{w} \cdot \mathbf{x})\end{aligned}$$
$$|f(\mathbf{x})| = |\mathbf{w} \cdot \mathbf{x}|$$

Loss Functions

- **Natural Loss:**

- Zero-One loss

$$\ell(\hat{y}, y) = \begin{cases} 0 & y = \hat{y} \\ 1 & y \neq \hat{y} \end{cases}$$


- **Real-valued-predictions loss:**

- Hinge loss

$$\ell(\hat{y}, y) = \max\{0, 1 - yf(\mathbf{x})\}$$

- Exponential loss (Boosting)

Loss Functions



Online Framework

- Initialize Classifier $f_1(\mathbf{x})$
- Algorithm works in rounds $t = 1 \dots T \dots$
- On round t the online algorithm :
 - Receives an input instance \mathbf{x}_t
 - Outputs a prediction $f_t(\mathbf{x}_t) = \hat{y}_t$
 - Receives a feedback label y_t
 - Computes loss $\ell(\hat{y}_t, y_t)$
 - Updates the prediction rule $f_t \rightarrow f_{t+1}$
- Goal :
 - Suffer small cumulative loss $\sum_t \ell(\hat{y}_t, y_t)$

Margin

- **Margin** of an example (\mathbf{x}_t, y_t) with respect to the classifier \mathbf{w}_t :


$$y_t(\mathbf{w}_t \cdot \mathbf{x}_t)$$

- Note : $y_t(\mathbf{w}_t \cdot \mathbf{x}_t) > 0 \quad \Leftrightarrow \quad \ell_{01}(\hat{y}_t, \mathbf{x}_t) = 0$


- The set $(\mathbf{x}_1, y_1) \dots (\mathbf{x}_T, y_T)$ is separable iff there exists \mathbf{u} such that

$$y_t(\mathbf{u} \cdot \mathbf{x}_t) > 0 \quad \forall t$$

Geometrical Interpretation



Hinge Loss



Why Online Learning?

- Fast
- Memory efficient - process one example at a time
- Simple to implement
- Formal guarantees – Mistake bounds
- Online to Batch conversions
- No statistical assumptions
- Adaptive

Update Rules

- Online algorithms are based on an update rule which defines f_{t+1} from f_t (and possibly other information)
- **Linear Classifiers** : find \mathbf{w}_{t+1} from \mathbf{w}_t based on the input (\mathbf{x}_t, y_t)
- **Some Update Rules** :
 - Perceptron (Rosenblatt)
 - ALMA (Gentile)
 - ROMMA (Li & Long)
 - NORMA (Kivinen et. al)
 - MIRA (Crammer & Singer)
 - EG (Littlestone and Warmuth)
 - Bregman Based (Warmuth)
 - CWL (Dredze et. al)

Design Principles of Algorithms

- If the learner suffers non-zero loss at any round, then we want to balance two goals:
 - **Corrective:** Change weights so that we **don't make this error again**
 - **Conservative:** **Don't change the weights too much**

$$d(w_{t+1}, w_t) + \eta L(y_t, w_{tt})$$

The Perceptron Algorithm

- If the learner suffers non-zero loss at any round, then we want to balance two goals:

$$d(w_{t+1}, w_t) + \eta L(y_t, w_{tt})$$

- **Euclidean distance**
- **Hinge loss**
- **Stochastic Gradient Descent (SGD)**

The Perceptron Algorithm ($\eta = 1$)

- If No-Mistake $y_t(\mathbf{w}_t \cdot \mathbf{x}_t) > 0$

- Do nothing $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t$

- If Mistake $y_t(\mathbf{w}_t \cdot \mathbf{x}_t) \leq 0$


- Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_t \mathbf{x}_t$

The Perceptron Algorithm ($\eta = 1$)

- When mistake happens, what does the update do?
 - **If $y_t = 1$:** $w_{t+1} = w_t + x_t$
 - ✓ w_{t+1} moves “closer to” x_t OR
 - ✓ x_t moves towards the positive side of the decision boundary
 - **If $y_t = -1$:** $w_{t+1} = w_t - x_t$
 - ✓ w_{t+1} moves “away from” x_t OR
 - ✓ x_t moves towards the negative side of the decision boundary
- In both cases, we are moving towards the “correct solution”

Running Example #1

Training Data: $((4, 0), 1), ((1, 1), -1), ((0, 1), -1), ((-2, -2), 1)$



Running Example #1

Round 1:

* $w_1 = 0$

Training Data:


$((4, 0), 1), ((1, 1), -1), ((0, 1), -1), ((-2, -2), 1)$

* $y_t \langle w_t, x_t \rangle$ for $t=1$

= 0 as well.

* $w_2 = w_1 + y_1 x_1$

= $(4, 0)$



Round 2:

* $y_2 \langle w_2, x_2 \rangle < 0$

* So $w_3 = w_2 + y_2 x_2$
 $= (4, 0) - (1, 1) = (3, -1)$


Round 3:

* $y_3 \langle w_3, x_3 \rangle > 0$

Correct. $w_4 = w_3 = (3, -1)$

Running Example #1

Training Data: $((4, 0), 1), ((1, 1), -1), ((0, 1), -1), ((-2, -2), 1)$



Round 4:

$$* y_4 \langle w_4, x_4 \rangle < 0$$

$$\begin{aligned} * \text{ So } w_5 &= w_4 + y_4 x_4 \\ &= (3, -1) + (-2, -2) \\ &= (1, -3) \end{aligned}$$

Running Example #2

$((1, 1), 1), ((1, -1), -1), ((-1, 1), -1), ((-1, -1), 1)$

Initially, $w_1 = 0$. In round 1,


* $y_1 \langle w_1, x_1 \rangle \leq 0$, so mistake.

* $w_2 = w_1 + y_1 x_1 = (1, 1)$


Round 2:

* $y_2 \langle w_2, x_2 \rangle \leq 0$

* $w_3 = w_2 + y_2 x_2 = (1, 1) - (1, -1)$
 $= (0, 2)$



Running Example #2



Round 3:


$$* y_3 \langle w_3, x_3 \rangle \leq 0$$

$$* w_4 = w_3 + y_3 x_3 = (0, 2) + (-1, 1) = (1, 1)$$

Round 4:

$$* y_4 \langle w_4, x_4 \rangle \leq 0.$$

$$* w_5 = w_4 + y_4 x_4 = (1, 1) + (-1, -1) = 0$$



The Perceptron Algorithm

- Suppose w_t makes a mistake on (x_t, y_t) , and we update w_{t+1} as $w_{t+1} = w_t + y_t x_t$. Is it possible for w_{t+1} to also make a mistake on (x_t, y_t) ?



The Perceptron Algorithm ($\eta = 1$)

- Suppose w_t makes a mistake on (x_t, y_t) , and we update w_{t+1} as $w_{t+1} = w_t + y_t x_t$. Is it possible for w_{t+1} to also make a mistake on (x_t, y_t) ?




- ▲ Yes, depends on the learning rate η


When does Perceptron converge?

- **Linear Separability**

- ▲ There exists a hyper-plane (weight vector) separating the positive and negative points



Linearly separable



Not linearly separable

Measure of Separability


- **Margin**

- ▲ For a weight vector \mathbf{w} , and training set \mathcal{S} , margin of \mathbf{w} with respect to \mathcal{S} is defined as follows:


$$\gamma(\mathbf{w}) = \min_{(x,y) \in \mathcal{S}} y(\mathbf{w} \cdot \mathbf{x})$$

- The training data \mathcal{S} is linearly separable if there exists *at least* one weight vector \mathbf{w} for which the margin is positive, i.e., $\gamma(\mathbf{w}) > 0$.

Margin: Examples



Low margin data



High margin data

Perceptron: Convergence Result

- **Theorem:** If the training data is linearly separable with margin γ , and if $\|x_i\| \leq 1$ for all examples (x_i, y_i) in the training set, then perceptron makes $\leq \frac{1}{\gamma^2}$ mistakes.
 - ▲ Proof??
- Lower margin implies more mistakes
- May need more than one pass over the training data to get a classifier with no mistakes

What if data is not linearly separable?

- Ideally, we want to find a linear separator that makes the minimum number of mistakes on the training data
 - ▲ NP-Hard problem! (Minsky and Papert, 1969)
 - ▲ This result killed the neural networks research in 1970's
- Perceptron still works
 - ▲ there will be few mistakes close to the decision boundary
 - ▲ will never converge to a single w as we make more passes

Problems with Perceptron

- Doesn't converge with inseparable data
 - Weight updates may often be very “bold”
 - Doesn't optimize margin
 - Sensitive to the order of examples
- ▲ **Voted and Averaged perceptron**

Voted Perceptron

- **Initialization:** $m = 1$; $w_1 = 0$; $c_m = 1$
- **Training Examples:** for $t = 1, 2, 3, \dots$
 - ▲ If mistake, update weights
 - $w_{m+1} = w_m + y_t x_t$
 - $m = m + 1$
 - $c_m = 1$
 - ▲ Else
 - $c_m = c_m + 1$ // counting how long w_m survived
- **Output:** $(w_1, c_1), (w_2, c_2), (w_3, c_3), \dots$

Voted Perceptron Classifier

$$f(x) = \textit{sign} \left(\sum_{i=1}^m c_i \textit{sign}(\langle w_i, x \rangle) \right)$$

- Any drawbacks of voted perceptron?

Voted Perceptron Classifier

$$f(x) = \text{sign} \left(\sum_{i=1}^m c_i \text{sign}(< w_i, x >) \right)$$

- Any drawbacks of voted perceptron?
- Yes, we have to store all the classifiers (in practice could be many)
- How can we solve this problem?

Averaged Perceptron

- Same algorithm as voted perceptron, but the classification rule is different

$$f_{average}(x) = \text{sign} \left(\sum_{i=1}^m (< c_i w_i, x >) \right)$$

$$f_{voted}(x) = \text{sign} \left(\sum_{i=1}^m c_i \text{sign}(< w_i, x >) \right)$$

Averaged vs. Voted Perceptron

- Simple Example: If $c_1 = c_2 = c_3 = 1$

$$f_{average}(x) = \textit{sign}(\langle w_1 + w_2 + w_3, x \rangle)$$

$$f_{voted}(x) = \textit{majority sign of } \langle w_1, x \rangle, \langle w_2, x \rangle, \langle w_3, x \rangle$$

Some Practical Tricks

- **Shuffling**

- ▶ shuffling the training examples in each iteration

- **Variable learning rate**

- ▶ decrease as learning progresses
- ▶ follow some schedule
- ▶ Set automatically by line search (converges faster)
- ▶ See Leon Bottou's SGD website:
<http://leon.bottou.org/projects/sgd>

- **Averaged Perceptron can be implemented very efficiently** (See Algorithm 7 in Hal's chapter)

Perceptron vs. Averaged Perceptron

- Perceptron

- ▲ $w_{t+1} = w_t + \eta \cdot y_t \cdot x_t$
- ▲ X_t has “n” features, but only “d” features are non-zero $\Rightarrow X_t$ is sparse
- ▲ $O(n) \Rightarrow O(d)$ If $d \ll n$, then it is a huge saving

- Averaged Perceptron

- ▲ $w = 0, w_{\text{sum}} = 0, \text{count} = 0$
- ▲ If we make a mistake, then we do the following:
 - Update $w = w + \delta$
 - Update $w_{\text{sum}} = w_{\text{sum}} + w \Rightarrow O(n)$
 - Update $\text{count} = \text{count} + 1$
- ▲ In the end, $w_{\text{avg}} = w_{\text{sum}} / \text{count}$

Some Practical Tricks

- **Learning Curve**

- ▶ Training iterations vs. number of mistakes
- ▶ You want to see that the mistakes decrease as we increase the no. of iterations (curve goes down)
- ▶ Very useful in debugging and seeing the behavior of online learning algorithms


- **Hyper-parameter Optimization**

- ▶ Split the training data: sub-train + validation data
- ▶ Tune hyper-parameters (e.g., no. of iterations) on the validation data
- ▶ The learner should not look at the test data!

Two kinds of learning curves

- General learning curve
 - ▲ Applicable for any learning algorithm
 - ▲ X-axis (number of training examples)
 - ▲ Y-axis (accuracy on unseen data)
- Online learning curve
 - ▲ Specific to online learning algorithms
 - ▲ For a fixed training set
 - ▲ X-axis (number of iterations or passes over the data)
 - ▲ Y-axis (number of mistakes made)

History of Perceptron*



* slide from Liang Huang

Passive-Aggressive (PA) Algorithm


PA Algorithm: Motivation

- **Perceptron:** No guaranties of margin *after* the update
- **PA:** Enforce a minimal non-zero margin after the update
- In particular :
 - ▲ If the margin is large enough (1), then do nothing
 - ▲ If the margin is less then unit, update such that the margin *after* the update is *enforced* to be unit

Margin

- $Y.(w_{t+1}.x) > 0$ then the prediction is correct
 - ▲ $+1 (> 0) > 0 \Rightarrow \text{Sign}(w_{t+1}.x) = +1$
 - ▲ $-1 (< 0) > 0 \Rightarrow \text{Sign}(w_{t+1}.x) = -1$
- If I'm making a mistake, then this margin is < 0
- If margin > 0 implies corrective behavior
 - ▲ Margin > 1 also implies corrective behavior

Input Space



Input Space vs. Version Space

- **Input Space :**

- ▶ Points are input data $y_t \mathbf{X}_t$
- ▶ One constraint is induced by weight vector \mathbf{W}
- ▶ Primal space
- ▶ Half space = all input examples that are classified correctly by a given predictor (weight vector)


$$\{y\mathbf{x} : \mathbf{w} \cdot (y\mathbf{x}) \geq 0\}$$

- **Version Space :**


- ▶ Points are weight vectors \mathbf{W}
- ▶ One constraints is induced by input data $y_t \mathbf{X}_t$
- ▶ Dual space
- ▶ Half space = all predictors (weight vectors) that classify correctly a given input example

$$\{\mathbf{w} : \mathbf{w} \cdot (y\mathbf{x}) \geq 0\}$$


Weight vector (Version) Space



Passive Step



Aggressive Step



The algorithm
projects w_t
on the desired
half-space

Aggressive Update Step

- Set \mathbf{w}_{t+1} to be the solution of the following optimization problem :

$$\begin{aligned} \mathbf{w}_{t+1} = \min_{\mathbf{w}} \quad & \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 && \text{Conservative} \\ \text{s.t.} \quad & y_t(\mathbf{w} \cdot \mathbf{x}_t) \geq 1 && \text{Corrective} \end{aligned}$$

- The Lagrangian :

$$\mathcal{L}(\mathbf{w}, \tau) = \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 + \tau(1 - y_t(\mathbf{w} \cdot \mathbf{x}_t))$$

- Solve for the dual : $\max_{\tau \geq 0} \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}, \tau)$

Norm of a vector

- L_1 norm
- L_2 norm
- L_infinity norm
- $W = [w_1, w_2, \dots, w_d]$
- $||W|| = \text{Sqrt}(w_1^2 + w_2^2 + \dots + w_d^2)$

Aggressive Update Step

- Optimize for \mathbf{w} :

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = \mathbf{w} - \mathbf{w}_t - \tau y_t \mathbf{x}_t$$

- Set the derivative to zero $\mathbf{w} = \mathbf{w}_t + \tau y_t \mathbf{x}_t$
- Substitute back into the Lagrangian :

$$\mathcal{L}(\tau) = -\frac{1}{2} \|\mathbf{x}_t\|^2 \tau^2 + \tau(1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t))$$

- Dual optimization problem

$$\max_{\tau \geq 0} \quad -\frac{1}{2} \|\mathbf{x}_t\|^2 \tau^2 + \tau(1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t))$$

Aggressive Update Step

- Dual Problem :

$$\max_{\tau \geq 0} -\frac{1}{2}\|\mathbf{x}_t\|^2\tau^2 + \tau(1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t))$$

- Solve it :

$$\tau = \max \left\{ 0, \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2} \right\}$$

Alternative Derivation

- Additional Constraint (linear update) :

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \tau y_t \mathbf{x}_t$$

- Force the constraint to hold as equality

$$1 = y_t((\mathbf{w}_t + \tau y_t \mathbf{x}_t) \cdot \mathbf{x}_t) = y_t(\mathbf{w}_t \cdot \mathbf{x}_t) + \tau \|\mathbf{x}_t\|^2$$

- Solve :

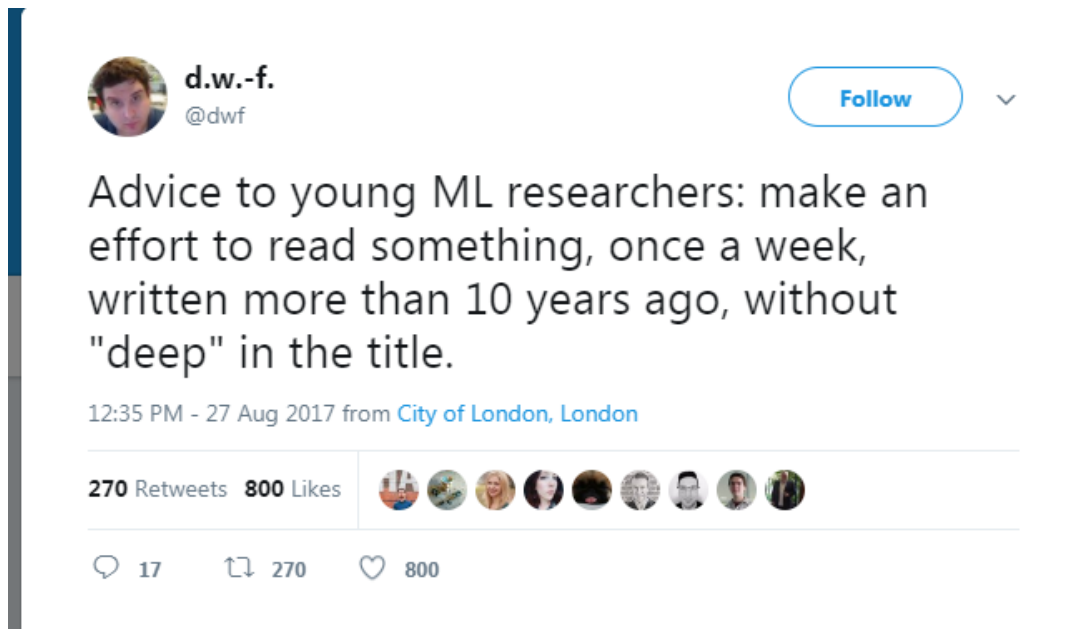
$$\tau = \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2}$$

Recap of last Lecture


- Perceptron algorithm
 - ▲ Convergence result for linearly separable data
 - ▲ Voted and Averaged perceptron
 - ▲ Practical tricks: shuffling, variable learning rate, ...
- History of Perceptron
 - ▲

Recap of Last Lecture

- **Passive-Aggressive Algorithm**
 - ▲ Smallest change to weights to ensure a margin of 1 on the new example
 - ▲ Derivation of weight update
 - ▲ Perceptron vs. Passive-Aggressive update



Passive-Aggressive Update



$$\mathbf{w}_{t+1} = \mathbf{w}_t + \tau y_t \mathbf{x}_t$$

$$y_t(\mathbf{w}_t \cdot \mathbf{x}_t) \geq 1$$

$$y_t(\mathbf{w}_t \cdot \mathbf{x}_t) < 1$$

$$\tau = 0$$

$$\tau = \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2}$$

Perceptron vs. PA Update

- Common Update :

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \tau y_t \mathbf{x}_t$$

- Perceptron

$$\tau = \begin{cases} 1 & y_t(\mathbf{w}_t \cdot \mathbf{x}_t) > 0 \\ 0 & y_t(\mathbf{w}_t \cdot \mathbf{x}_t) \leq 0 \end{cases}$$

- Passive-Aggressive

$$\tau = \max \left\{ 0, \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2} \right\}$$

Two kinds of updates


- Additive update
 - ▲ $w_{t+1} = w_t + \Delta$ (adding some gradient)
- Multiplicative update
 - ▲ $w_{t+1} = w_t \times \Delta$
 - ▲ Example: Exponentiated Gradient Algorithm
- Yoram Singer => Hebrew Univ => Google (2007)
 - ▲ Spam classifiers in Gmail => they are passive-aggressive algorithm

The Passive-Aggressive Algorithm


- Each example defines a set of consistent hypotheses: $C_\varepsilon(\mathbf{z}_t) = \{\mathbf{w} \mid \delta(\mathbf{w}; \mathbf{z}_t) \leq \varepsilon\}$
- The new vector \mathbf{w}_{t+1} is set to be the projection of \mathbf{w}_t onto $C_\varepsilon(\mathbf{z}_t)$

$$\mathbf{w}_{t+1} = \underset{\mathbf{w}}{\operatorname{argmin}} \|\mathbf{w} - \mathbf{w}_t\| \quad \text{s.t.} \quad \mathbf{w} \in C_\varepsilon(\mathbf{z}_t)$$

Classification



Regression



Unrealizable Case



There is no weight vector that satisfy all the constraints

Unrealizable Case

$$\mathbf{w}_{t+1} = \min_{\mathbf{w}} \quad \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 + C\xi^2$$

s.t. $y_t(\mathbf{w} \cdot \mathbf{x}_t) \geq 1 - \xi \quad \xi \geq 0$



$$\mathbf{w}_{t+1} = \mathbf{w}_t + \tau y_t \mathbf{x}_t$$



$$\min \left\{ C, \max \left\{ 0, \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2} \right\} \right\}$$



$$\max \left\{ 0, \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2 + C} \right\}$$

Reduction Hammer


- Smaller problem/nail (binary classification) => Simpler solution/hammer (Perceptron or PA)
 - ▲ You have only two classes
- Harder problem/nail
 - ▲ You have more than two classes
 - ▲ How can I develop an appropriate hammer for this nail?
- Create some simpler problems that you know how to solve => Solve them => Aggregate the solutions to come up with a solution for the harder problem

Reduction for Multi-Class Classification


- Classes = {red, blue, green}
- Bunch of training examples $\{x, y\}$ pairs where y is one of the three classes
- One against others {red vs. other; blue vs. others; green vs. others}
- Increasing the dimension of margin
- Tree of classifiers to filter candidate class labels
- One vs. One: Red vs. Blue; Blue vs. Green; Green vs. Red => Given a new example => Majority vote

Multi-Class: Representation-I

- k Prototypes $w_1, w_2 \dots w_k$
- New instance x
- Compute $\text{Score}(r) = w_r \cdot x$



Class r	$w_r \cdot x$
1	-1.08
2	1.66
3	0.37
4	-2.09



- Prediction:
The class achieving the highest Score

Multi-Class Representation-II

- Weight-vector per class (Representation I)
 - ▲ Intuitive
- Single weight-vector (Representation II)
 - ▲ Generalizes representation I

$$F(x,4) = \begin{array}{|c|c|c|c|c|} \hline 0 & 0 & 0 & x & 0 \\ \hline \end{array}$$

- Predict label with highest score (Inference)

$$\arg \max_z F(x, z) \cdot w$$

Margin for Multi-Class

- Binary :

$$\mathbf{w}_2 \cdot \mathbf{x} - \mathbf{w}_1 \cdot \mathbf{x} \geq 1$$

$$\mathbf{w} \cdot F(\mathbf{x}, 2) - \mathbf{w} \cdot F(\mathbf{x}, 1) \geq 1$$

- Multi Class :

$$\mathbf{w}_y \cdot \mathbf{x} - \mathbf{w}_z \cdot \mathbf{x} \geq 1 \quad \forall z \neq y$$

$$\mathbf{w} \cdot F(\mathbf{x}, y) - \mathbf{w} \cdot F(\mathbf{x}, z) \geq 1$$

$\forall z \neq y$

Margin for Multi-Class

- Multi Class :

$$w \cdot F(x, y) - w \cdot F(x, z) \geq 1$$

$\forall z \neq y$

Because the loss
function is not
constant !



Margin for Multi-Class


- Multi Class :

$$\mathbf{w} \cdot F(\mathbf{x}, y) - \mathbf{w} \cdot F(\mathbf{x}, z) \geq \ell(z, y) \\ \forall z \neq y$$


So, use it !



Margin Scaled by Loss: Illustration



Margin Scaled by Loss: Illustration



PA Multi-Class Update

- Project the current weight vector such that the instance ranking is consistent with loss function
- Set \mathbf{w}_{t+1} to be the solution of the following optimization problem :

$$\begin{aligned} \mathbf{w}_{t+1} &= \min_{\mathbf{w}} \quad \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 \\ \text{s.t.} \quad &\mathbf{w} \cdot F(\mathbf{x}_t, y_t) - \mathbf{w} \cdot F(\mathbf{x}_t, z) \geq \ell(z, y_t) \\ &\forall z \neq y_t \end{aligned}$$

PA Multi-Class Update

- **Problem**

- ▶ intersection of constraints may be empty

$$\left\{ \mathbf{w} : \begin{array}{l} \mathbf{w} \cdot F(\mathbf{x}_t, y_t) - \mathbf{w} \cdot F(\mathbf{x}_t, z) \geq \ell(z, y_t) \\ \forall z \neq y_t \end{array} \right\} = \emptyset$$

- **Solutions**

- ▶ Does not occur in practice
- ▶ Add a slack variable
- ▶ Remove constraints

Add a Slack Variable

- Add a slack variable :

$$\begin{aligned} \mathbf{w}_{t+1} &= \min_{\mathbf{w}} \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 + C\xi \\ \text{s.t.} \quad &\mathbf{w} \cdot F(\mathbf{x}_t, y_t) - \mathbf{w} \cdot F(\mathbf{x}_t, z) \geq \ell(z, y_t) - \xi \\ &\forall z \neq y_t \\ &\xi \geq 0 \end{aligned}$$

- Rewrite the optimization :

$$\frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 + C \max \left\{ \begin{array}{c} \mathbf{w} \cdot F(\mathbf{x}_t, z) - \mathbf{w} \cdot F(\mathbf{x}_t, y_t) + \ell(z, y_t) \\ 0 \end{array} \right\}$$

PA Multi-Class Update

- Remove constraints :

$$\begin{aligned} \mathbf{w}_{t+1} &= \min_{\mathbf{w}} \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 \\ \text{s.t.} \quad & \mathbf{w} \cdot F(\mathbf{x}_t, y_t) - \mathbf{w} \cdot F(\mathbf{x}_t, z) \geq \ell(z, y_t) \\ & \forall z \neq y_t \end{aligned}$$

$$\begin{aligned} \mathbf{w}_{t+1} &= \min_{\mathbf{w}} \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 \\ \text{s.t.} \quad & \mathbf{w} \cdot F(\mathbf{x}_t, y_t) - \mathbf{w} \cdot F(\mathbf{x}_t, \hat{y}_t) \geq \ell(\hat{y}_t, y_t) \end{aligned}$$

- How to choose the single competing labeling?
 - The labeling that attains the highest score!

$$\hat{y}_t = \arg \max_z \mathbf{w}_t \cdot F(\mathbf{x}_t, z)$$

- ... which is the predicted label according to the current model

PA Multi-Class Online Algorithm

- Initialize \mathbf{w}_1
- For $t = 1 \dots T \dots$
 - Receive an input instance \mathbf{x}_t
 - Outputs a prediction $\hat{y}_t = \arg \max_z \mathbf{w}_t \cdot F(\mathbf{x}_t, z)$
 - Receives a feedback label y_t
 - Computes loss $\ell(\hat{y}_t, y_t)$
 - Update the prediction rule

$$\begin{aligned} \mathbf{w}_{t+1} &= \min_{\mathbf{w}} \frac{1}{2} \|\mathbf{w} - \mathbf{w}_t\|^2 \\ \text{s.t.} \quad &\mathbf{w} \cdot F(\mathbf{x}_t, y_t) - \mathbf{w} \cdot F(\mathbf{x}_t, \hat{y}_t) \geq \ell(\hat{y}_t, y_t) \end{aligned}$$

Multi-Class Online Algorithm

- Initialize \mathbf{w}_1
- For $t = 1 \dots T \dots$
 - Receive an input instance \mathbf{x}_t
 - Outputs a prediction $\hat{y}_t = \arg \max_z \mathbf{w}_t \cdot F(\mathbf{x}_t, z)$
 - Receives a feedback label y_t
 - Computes loss $\ell(\hat{y}_t, y_t)$
 - Update the prediction rule

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \tau \cdot (F(\mathbf{x}_t, y_t) - F(\mathbf{x}_t, \hat{y}_t))$$

Recap of Last Lecture

- **Multi-Class Representations**

- ▶ Multi-prototype (one weight vector \mathbf{w}_i for each class i)
- ▶ Single-prototype (one weight vector \mathbf{w} – concatenation of all the k weight vector) via extended feature space $F(\mathbf{x}, \mathbf{y})$

- **Multi-Class Passive-Aggressive Algorithm**

- ▶ Margin for multi-class classification
- ▶ Margin scaled by the loss function
- ▶ Mathematical optimization problem with $k - 1$ constraints

Multi-Class Online Algorithm

- Initialize \mathbf{w}_1
- For $t = 1 \dots T \dots$
 - Receive an input instance \mathbf{x}_t
 - Outputs a prediction $\hat{y}_t = \arg \max_z \mathbf{w}_t \cdot F(\mathbf{x}_t, z)$
 - Receives a feedback label y_t
 - Computes loss $\ell(\hat{y}_t, y_t)$
 - Update the prediction rule

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \tau \cdot (F(\mathbf{x}_t, y_t) - F(\mathbf{x}_t, \hat{y}_t))$$

$$\tau = \frac{1 - y_t(\mathbf{w}_t \cdot \mathbf{x}_t)}{\|\mathbf{x}_t\|^2}$$

Binary

$$\tau = \frac{1 - (\mathbf{w}_t \cdot F(\mathbf{x}_t, y_t) - \mathbf{w}_t \cdot F(\mathbf{x}_t, \hat{y}_t))}{\|F(\mathbf{x}_t, y_t) - F(\mathbf{x}_t, \hat{y}_t)\|^2}$$

Multi-class

Multi-Class Classification Implementation

- “d” features for each input example
- W_1 (d weights) for class 1
- W_2 (d weights) for class 2
- ..
- $\text{Score}(\text{class } 1) = W_1 \cdot X$
- $W_{\{\text{correct}\}} = W_{\{\text{correct}\}} + \eta \cdot X$
- $W_{\{\text{wrong}\}} = W_{\{\text{wrong}\}} - \eta \cdot X$

Multi-Class Update

- $\text{Score}(\text{correct-label}) = 3.4$
 - ▲ $W_{\{\text{correct-label}\}}.X$
- $\text{Score}(\text{wrong-label}) = 5.6$
 - ▲ $W_{\{\text{wrong-label}\}}.X$
- Weight update
 - ▲ $W_{\{\text{correct-label}\}} = W_{\{\text{correct-label}\}} + \tau.X$
 - ▲ $W_{\{\text{wrong-label}\}} = W_{\{\text{wrong-label}\}} - \tau.X$
- Updated score(correct-label)
 - ▲ $(W_{\{\text{correct-label}\}} + \tau.X).X = \tau.X^2$ (added)
- Updated score(wrong-label)

Binary vs. Multi-Class Interpretation

- Correct-label y_t
- Wrong-label \hat{y}_t
- Feature vectors
 - ▲ $F(x, y_t)$
 - ▲ $F(x, \hat{y}_t)$
- $F(x, y_t) - F(x, \hat{y}_t) == \text{Input feature vector}$
- Output label = +1
- $F(x, \hat{y}_t) - F(x, y_t) == \text{Input feature vector}$
- Output label = -1

Confusion Matrix

- Measures which classes are easy or hard to separate -- k classes implies $k \times k$ matrix

k

k

- $C_{ij} = \frac{\text{\# examples with class label } j \text{ that are classified as label } i}{\text{\# examples with label } j}$
- High diagonal entry implies class is easy to classify
- High off-diagonal entry implies classes are easily confused

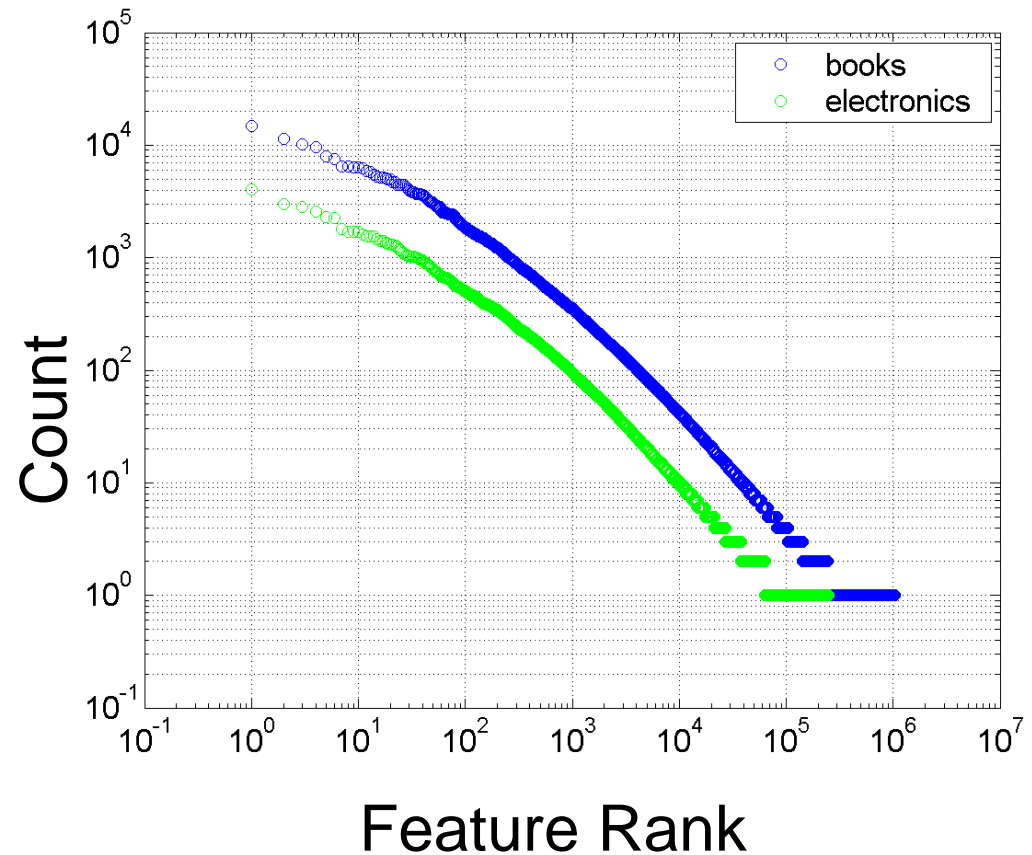
Learning Curves: Online vs. General

- **Online learning curve**
 - ▶ training iterations vs. Mistakes
 - ▶ only applicable for online learning algorithms
- **General learning curve**
 - ▶ amount of training data vs. accuracy
 - ▶ applicable to any learning algorithm


Confidence-Weighted (CW) Algorithm

Motivation: NLP problems

- Big datasets, large number of features
- Many features are only weakly correlated with target label
- Heavy-tailed feature distribution
- Linear classifiers: features are associated with word counts



Motivation: Sentiment Classification



- Who needs this Simpsons book?
You Doooooooooooo

This is one of the most extraordinary volumes I've ever encountered encapsulating a television series Exhaustive, informative, and ridiculously entertaining, it is the best accompaniment to the best television show Even if you only "enjoy" the Simpsons (as opposed to being a raving fanatic, which most people who watch the show are, after all ... Very highly recommended!



Motivation: Sentiment Classification

- Many positive reviews with the word **best**

$$W_{\text{best}}$$

- Later negative review

▲ “*boring book – best if you want to sleep in seconds*”

- Linear update will reduce both

$$W_{\text{best}} \quad W_{\text{boring}}$$

- But **best** appeared more than **boring**

- How to adjust weights at different rates? W_{boring} W_{best}

Span based Update Rules

- The weight vector is a linear combination of examples

$$w_f \leftarrow w_f + \eta y x_f \quad \forall f$$

Weight of
feature f

Learning rate

Target label, -
1 or 1

Value of feature
 f of instance

- Two rate schedules (among others):


- ▲ Perceptron algorithm, conservative:

$$\eta = 1$$

- ▲ Passive-aggressive

$$\eta = \max \left\{ 0, \frac{1 - y(\mathbf{w} \cdot \mathbf{x})}{\|\mathbf{x}\|^2} \right\}$$

Distributions in Version Space



Margin as a Random Variable

- Signed margin

$$M = y(w \cdot x)$$

is a Gaussian-distributed variable

$$M \sim \mathcal{N} \left(y(x \cdot \mu) , x^\top \Sigma x \right)$$

- Thus:

$$\Pr [y (w \cdot x) \geq 0] = \Phi \left(\frac{y(x \cdot \mu)}{\sqrt{x^\top \Sigma x}} \right)$$

$$\Phi(z) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^z e^{-t^2} dt$$

Version (weight vector) Space

Place most of the probability mass in this region

x_t

$$0 \leq y_t(w \cdot x_t)$$

Passive Step

Nothing to do, most weight vectors already classify the example correctly


μ_t, Σ_t

$$0 \leq y_t(w \cdot x_t)$$

Aggressive Step

Mean moved past
the mistake line
(large margin)

The covariance is
shirked in the
direction of the
new example



PA like Update

- PA:

$$\begin{array}{ll} \min_{\mathbf{w}} & \frac{1}{2} \|\mathbf{w} - \mathbf{w}_i\|^2 \\ \text{s.t.} & y_i(\mathbf{w} \cdot \mathbf{x}_i) \geq 1 \end{array}$$

- New Update :

$$\begin{array}{ll} \min_{\mu, \Sigma} & D_{\text{KL}}(\mathcal{N}(\mu, \Sigma) \parallel \mathcal{N}(\mu_i, \Sigma_i)) \\ \text{s.t.} & \Pr[y_i(\mathbf{w} \cdot \mathbf{x}_i) \geq 0] \geq \eta \end{array}$$

Confidence
Parameter

$$\eta \in (0.5, 1)$$

η

Summary of Online Learning

- **Online learning**
 - ▲ Iterative game between teacher and learner
- **Design principles of online learning**
 - ▲ Trade-off amount of change (conservative) and reduction in loss (corrective)
- **Online learning algorithms**
 - ▲ Perceptron (fixed learning rate for all examples)
 - ▲ Passive-Aggressive (fixed learning rate for each example)
 - ▲ Confidence-weighted classifier (fixed learning for each feature and each example)

Questions?