

Midterm Exam #2 Review


Janardhan Rao (Jana) Doppa

School of EECS, Washington State University


What is Machine Learning?

- Machine learning = Automating Automation

Traditional Programming



Machine Learning




Training data

Learning Paradigms


- **Supervised Learning – Good coverage**
- **Semi-Supervised Learning – No coverage**
- **Unsupervised Learning – Limited coverage**
- **Active Learning – High level overview**
- **Reinforcement Learning – Good coverage**

Supervised Learning for Simple Outputs

Training Data
 $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$



θ



Logistic Regression
Support Vector Machines
K Nearest Neighbor
Decision Trees
Neural Networks

Example problem:

X - image of a face

$Y \in \{\text{male, female}\}$

X

feature vector

Y

class label

Overview of ML Algorithms

- There are lot of machine learning algorithms
- Every machine learning algorithm has three components
 - ▲ **Representation**
 - ▲ **Evaluation**
 - ▲ **Optimization**

Online Learning

- **Online learning**
 - ▲ Iterative game between teacher and learner
- **Design principles of online learning**
 - ▲ Trade-off amount of change (conservative) and reduction in loss (corrective)
- **Online learning algorithms**
 - ▲ Perceptron (fixed learning rate for all examples)
 - ▲ Passive-Aggressive (fixed learning rate for each example)
 - ▲ Confidence-weighted classifier (fixed learning for each feature and each example)

Support Vector Machine (SVM) Classifier

- **Maximum margin classification:** Linear hyperplane that maximizes margin
- **Support vectors:** a small number of training examples are important to construct the classifier
- **Non-separable data**
 - ▲ **Soft-margin formulation:** relax the constraints
 - ▲ **Kernel-trick:** implicit mapping of data into high-dimensional space without any additional computational burden
- **Kernelizing online learning algorithms**
 - ▲ **Primal:** update weights directly
 - ▲ **Dual:** update the coefficients of the training examples

Non-Parametric Classifiers

- **K-nearest neighbor classification**
 - ▲ Flexible hypothesis: infinitely complex
 - ▲ Classification time is high (NN computation)
 - ▲ Doesn't work very well in high dimensions
- **Decision tree classification**
 - ▲ Hierarchical partitioning of the input space into axis parallel rectangles
 - ▲ ID3 decision tree construction: greedy procedure based on information gain heuristic
 - ▲ Prone to over-fitting, but can be mitigated by pruning based on the validation data

Probabilistic Classifiers

- **Logistic Regression classifier**

- ▶ Represent probability distribution as a parametric sigmoid function
- ▶ Learn parameters via maximum likelihood estimation
- ▶ Closely related to the Perceptron classifier

- **Naïve Bayes classifier**

- ▶ Learns $P(Y)$ and $P(X|Y)$
- ▶ **Assumption:** each feature is independent from one another given the class label
- ▶ Drastically reduces the number of parameters: improves computational and sample efficiency
- ▶ Lot of success in real-world applications (e.g., text classification)

Ensemble Methods

- **Meta Learning:** combine multiple classifiers into a single one to improve the performance
- **Bagging** (Bootstrap AGGREGatING)
 - ▲ Learn multiple classifiers by subsampling the training data
 - ▲ Classify new examples via majority vote
- **Boosting**
 - ▲ Combining multiple simple rules to come up with a highly accurate rule (or classifier)
 - ▲ Iteratively modify the weights of the training examples based on their hardness (i.e., misclassified by previous rules of thumb)
 - ▲ Classify new examples via weighted majority of the rules
 - ▲ **AdaBoost:** a concrete algorithm for boosting

Bias and Variance Decomposition

- **Under-fitting:** Models with too few parameters can perform poorly
- **Over-fitting:** Models with too many parameters can perform poorly
- Generalization error can be decomposed into Bias and Variance
 - ▲ Simple (inflexible) models will have high bias and Complex (flexible) models will have high variance
 - ▲ **Bias:** measures the accuracy or quality of the algorithm
 - ▲ **Variance:** measures the precision or specificity of match
 - ▲ **Tradeoff:** low bias \Rightarrow high variance and vice versa
- Bagging reduces variance and Boosting reduces bias

Applying ML in Practice

- **Over-fitting**

- ▶ Training accuracy is much higher than the testing accuracy
- ▶ Model explains the training set very well, but poor generalization
- ▶ More training data and less features will help

- **Under-fitting**

- ▶ Both training and testing accuracies are very low
- ▶ Model cannot represent the target concept well enough
- ▶ More features and using an expressive model will help

- Error analysis and ablation analysis will help debug and improve end-to-end systems
- Ensembles (e.g., random forests) works really well
- Learning == Generalization

Unsupervised Learning

- **Setting:** we are only provided with examples without specifying the labels and we want to cluster the data into groups
- Define a **distance measure** between input examples
 - ▲ Symmetric, positivity and self-similarity, satisfy triangle inequality
- **Hierarchical clustering algorithms**
 - ▲ Bottom up (agglomerative) and top down (divisive)
- **Partition clustering algorithms**
 - ▲ K-Means and Mixture of Gaussians (model based clustering)
 - ▲ Expectation Maximization (EM) algorithm for learning from hidden (missing) data
 - ▲ K-Means == Hard EM; and GMMs == Soft EM

Reinforcement Learning (1)

- **Problem:** Learning to Act (take decisions) by interacting with a system (world) to maximize the cumulative reward
- World is modeled as a **Markov Decision Process (MDP)**
 - ▲ Finite states, finite actions, stochastic transition function, and bounded real-valued reward function
- **Assumptions**
 - ▲ First-order Markovian dynamics
 - ▲ State-dependent reward
 - ▲ Stationary dynamics
 - ▲ Full observability
- Solution: policies (“plans” for MDPs)

Reinforcement Learning (2)

- **Non-stationary policy**

- ▶ $\pi: S \times T \rightarrow A$; $\pi(s, t)$ tells us what action to take at state s when there are t stages-to-go
- ▶ Need when we are given a finite planning horizon H

- **Stationary policy**

- ▶ $\pi: S \rightarrow A$; $\pi(s)$ is action to do at state s (regardless of time)
- ▶ Need when we want to continue taking actions indefinitely

- **Value of a policy π at state s**

- ▶ Depends on immediate reward, but also what you achieve subsequently by following that policy

$$\begin{aligned} V_{\pi}^k(s) &= E \left[\sum_{t=0}^k R^t \mid \pi, s \right] \\ &= E \left[\sum_{t=0}^k R(s^t) \mid a^t = \pi(s^t, k-t), s^0 = s \right] \end{aligned}$$

RL Algorithms: Big Picture

Planning with **known** model (MDP)

- **Policy evaluation:**
 - ▶ Given an MDP and a (non)stationary policy π
 - ▶ Compute finite-horizon value function $V_{\pi}^k(s)$ for any k
- **Policy optimization:**
 - ▶ Given an MDP and a horizon H
 - ▶ Compute the optimal finite-horizon policy
 - ▶ Equivalent to computing optimal value function (value iteration)

Planning with **unknown** model (MDP)

- **Policy evaluation:**
 - ▶ Given a stationary policy π , compute the value of policy
 - ▶ **Passive RL**: direct estimation, ADP, TD methods
- **Policy optimization:**
 - ▶ Compute the optimal policy
 - ▶ **Active RL** — ADP, TD, and Q learning

Finite-Horizon: Policy Evaluation


- Can use dynamic programming to compute $V_{\pi}^k(s)$
 - ▲ Markov property is critical for this

$$(k=0) \quad V_{\pi}^0(s) = R(s), \quad \forall s$$

$$(k>0) \quad V_{\pi}^k(s) = R(s) + \underbrace{\sum_{s'} T(s, \pi(s, k), s') \cdot V_{\pi}^{k-1}(s')}_{\text{expected future payoff with } k-1 \text{ stages to go}}, \quad \forall s$$

immediate reward

expected future payoff
with $k-1$ stages to go



Finite Horizon: Policy Optimization

- Markov property allows exploitation of DP principle for optimal policy construction
 - ▲ no need to enumerate $|A|^{Hn}$ possible policies

- Value Iteration

$$V^0(s) = R(s), \quad \forall s$$

Bellman backup

$$V^k(s) = R(s) + \max_a \sum_{s'} T(s, a, s') \cdot V^{k-1}(s')$$

$$\pi^*(s, k) = \arg \max_a \sum_{s'} T(s, a, s') \cdot V^{k-1}(s')$$

V^k is optimal k-stage-to-go value function

$\pi^*(s, k)$ is optimal k-stage-to-go policy

Passive RL: Policy Evaluation w/ unknown MDP

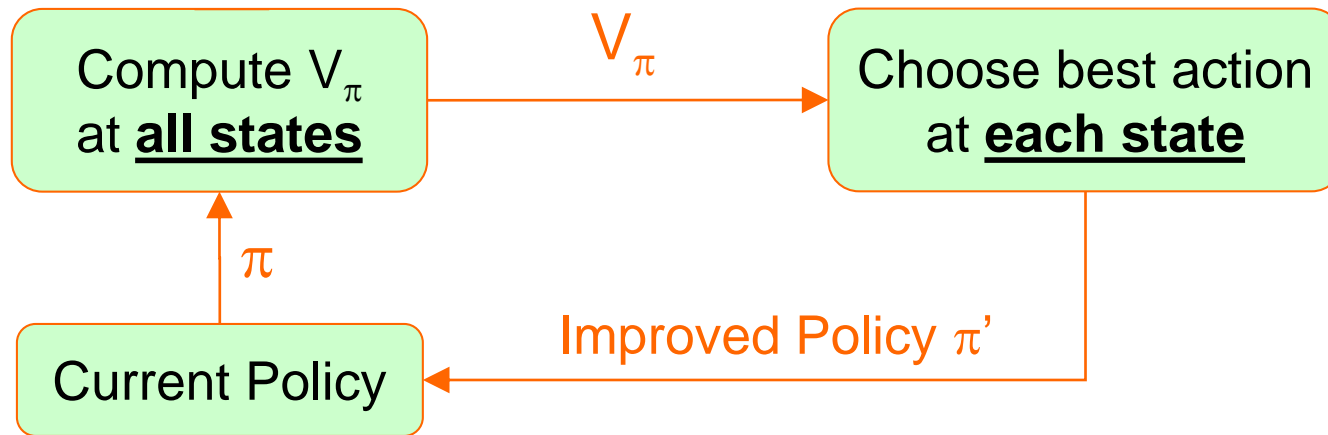
- Monte-Carlo Direct Estimation (model free)
 - ▲ Simple to implement
 - ▲ Each update is fast
 - ▲ Does not exploit Bellman constraints
 - ▲ Converges slowly
- Adaptive Dynamic Programming (model based)
 - ▲ Harder to implement
 - ▲ Each update is a full policy evaluation (expensive)
 - ▲ Fully exploits Bellman constraints
 - ▲ Fast convergence (in terms of updates)
- Temporal Difference Learning (model free)
 - ▲ Update speed and implementation similar to direct estimation
 - ▲ Partially exploits Bellman constraints---adjusts state to 'agree' with observed successor
 - Not **all** possible successors as in ADP
 - ▲ Convergence in between direct estimation and ADP

Active RL: Policy Optimization w/ unknown MDP

- **Exploration vs. Exploitation trade-off**
 - ▲ **Exploitation**: To try to get reward. We exploit our current knowledge to get a payoff.
 - ▲ **Exploration**: Get more information about the world. How do we know if there is not a pot of gold around the corner.
- **Basic intuition behind most approaches**
 - ▲ Explore more when knowledge is weak. Exploit more as we gain knowledge.
- **Exploration policy**
 - ▲ We want a policy that is greedy in the limit of infinite exploration (GLIE)
- **ADP-based (model based) RL**
 - ▲ Solve for optimal policy given the current model. Take action according to exploration policy. Update model based on new observation. Repeat.
- **TD-based (model based) RL**
 - ▲ Start with initial value function. Take action according to exploration policy. Update model based on new observation. Perform TD update to get new value function. Repeat.
- **Q-Learning (model free) RL**
 - ▲ Start with initial Q values. Take action according to exploration policy. Perform TD update to get new Q values. Repeat.

Approximate Policy Iteration for Large MDPs

Policy Iteration

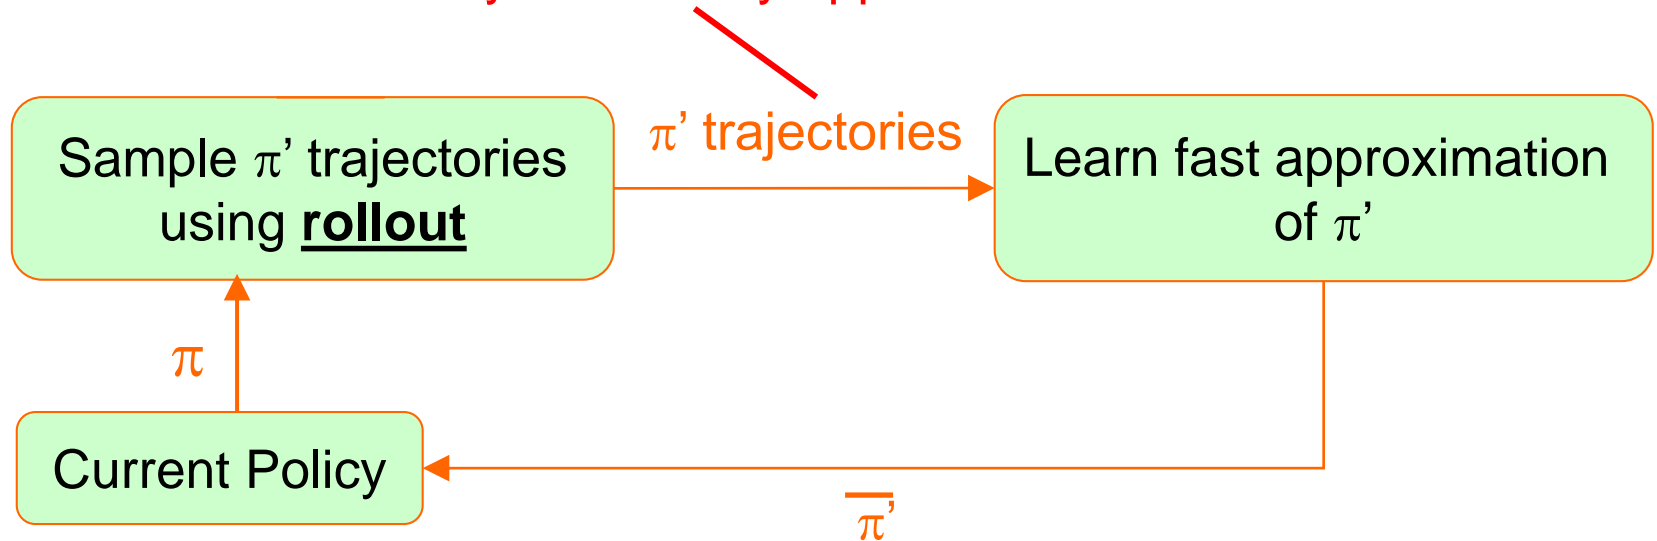


Approximate policy iteration:

- Only computes values and improved action at some states.
- Uses those to infer a fast, compact policy over all states.

Approximate Policy Iteration

technically rollout only approximates π' .



1. Generate trajectories of rollout policy (starting state of each trajectory is drawn from initial state distribution I)
2. “Learn a fast approximation” of rollout policy
3. Loop to step 1 using the learned policy as the base policy

Hyper-parameter Search via Bayesian Optimization

- **Build a surrogate statistical model** based on past computational experiments
 - ▲ Assumption is that **surrogate model is cheap to evaluate**
- **Intelligently select the next experiment** (candidate solution) **using the statistical model**
 - ▲ **Trade-off exploration and exploitation**
 - ▲ Exploration corresponds to selecting candidates for which the statistical model is not confident (high variance)
 - ▲ Exploitation corresponds to selecting candidates for which the statistical model is highly confident (high mean)

Hyper-parameter Search via Bayesian Optimization

- Initialize statistical model F
- Repeat the following steps for several iterations
 - ▲ Select the next candidate (say x) by optimizing the acquisition function $A(x)$
 - ▲ Run experiment with candidate x to compute its quality y
 - ▲ Update the statistical model F based on the new training example (x, y)
 - ▲ Update the best uncovered solution so far (say $x_{\{best\}}$)

Active Learning

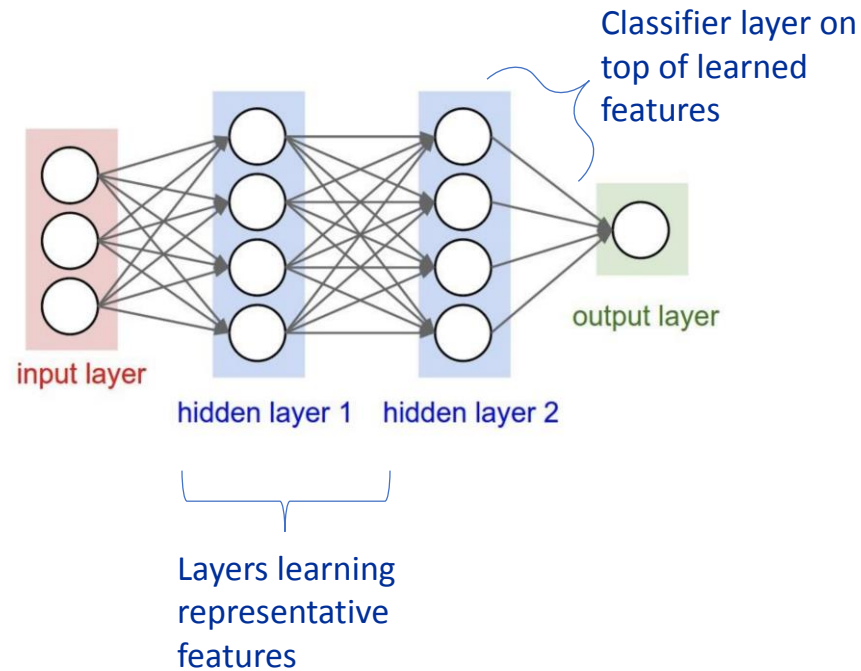
- Active learning is a **label-efficient** learning strategy
- Intelligently selects the examples based on their informativeness
- **Query Selection Strategies**
 - ▲ Uncertainty Sampling
 - ▲ Query By Committee (QBC)

Deep Learning: Differentiable Programming Paradigm

- The user writes a differentiable program
- Use training data to optimize the parameters of this program so that the program behaves as desired


Different ways to think about a neural network

Automatic Feature Learning



Different ways to think about a neural network

Function Approximator




$$x \longrightarrow f \longrightarrow y$$

$$f(x) = \sigma[W_2 [\sigma(W_1 x + b_1)] + b_2]$$

Inductive Biases and Architectures

- Similarly, Neural networks are a great tool!
- But they are very generic
- We need to add inductive biases in these models
- How do we do that?
- Different types of architectures!
- For e.g. Convolutional architecture for vision
- For e.g. Recurrent architecture for sequences, ...

This is what a CNN looks like



We will talk about each layer in detail shortly!

Learning Theory

- **Sample complexity:** How many training examples are needed for a learner to construct (with high probability) a highly accurate concept?
- **Computational complexity:** How much computational resources (time and space) are needed for a learner to construct (with high probability) a highly accurate concept?
- **PAC Model** (Leslie Valiant got a Turing Award!)
 - ▲ Only requires a **Probably Approximately Correct** (PAC) concept: learn a decent approximation most of the time
 - ▲ Requires polynomial sample complexity and computational complexity

PAC Learning

- The only reasonable expectation of a learner is that with *high probability* it learns a *close approximation* to the target concept
- In the PAC model, we specify two parameters, ϵ and δ , and require that with probability at least $(1 - \delta)$ a system learn a concept with error at most ϵ
- How to prove PAC learnability?
 - ▲ First, prove sample complexity of learning a *target concept h^** using a *hypothesis space H* is polynomial
 - ▲ Second, prove that the learner can train on a polynomial-sized data set in polynomial time